

```

1  /*
2  * Aaron Chan
3  * ECE 373 (Spring 2017)
4  * Assignment #6
5  *
6  * In assignment #4, we initialized a PCI driver
7  * to blink the LEDs on the Atom Box using a timer
8  * and adding ways to customize a blink rate to change
9  * the speed that it blinks.
10 *
11 * In this assignment, we will make the following changes:
12 * - small set of legacy descriptors for the receive queue (16 descriptors)
13   with buffers allocated to 2048 bytes
14   - control mechanism for keeping track of HEAD and TAIL of receive queue
15   - interrupt handler, tied to legacy interrupt source
16   - workqueue thread to handle deferred processing from interrupt
17   - configure code to put chip into promiscuous mode, and force link up at 1 Gb
18   - Update read() system call to return 16-bit unsigned int, lower 8 bits being
19     value received from receive queue TAIL, upper 8 bits being value from the HEAD.
20
21   As suggested, this assignment is done with MSI instead of Legacy descriptors.
22 */
23
24 #include <linux/module.h>
25 #include <linux/types.h>
26 #include <linux/kdev_t.h>
27 #include <linux/fs.h>
28 #include <linux/cdev.h>
29 #include <linux/slab.h>
30 #include <linux/uaccess.h>
31 #include <linux/pci.h>
32 #include <linux/timer.h>
33 #include <linux/delay.h>
34 #include <linux/workqueue.h>
35 #include <linux/interrupt.h>
36
37 // ===== MACROS =====
38 #define DEVCNT 5
39 #define DEVNAME "hw6_pci_interrupts"
40 #define DEV_NODE_NAME "hw6_interrupts"
41
42 // Device Info for Intel 82583v
43 #define VENDOR_ID 0x8086
44 #define DEVICE_ID 0x150c
45
46 #define LED_REG 0x00E00 // LED control register offset
47 #define LED_ON 0x4E4E0F //Green LEDs on value
48 #define LED_OFF 0x0F0F0F //All LEDs off value
49
50 // Chicken Bit necessities
51 #define GCR 0x5B00 // 3GIO Control Reg
52 #define GCR2 0x5B64 // 3GPIO Control Reg 2
53 #define MDIC 0x0020 // MDI Control Reg
54 #define STATUS 0x08
55
56 // Necessary for IRQ
57 #define ICR 0xC0 // Cause Read
58 #define ICS 0xC8 // Cause Setup
59 #define IMS 0xD0 // Mask Set
60 #define IMC 0xD8 // Mask Clear
61 #define CTRL 0x04 // Device Control Registering
62 #define RCTL 0x00100 // Receive Control Reg

```

```

63 #define IRQ_ENABLE 0x1100D4
64 #define PROMISCUOUS 0x801A
65 #define MAC 0x100A41
66
67 // Receive Descriptor
68 #define MAX 16 // Number of Descriptors
69 #define RDBAL 0x2800 // Base Address Low
70 #define RDBAH 0x2804 // Base Address High
71 #define RDLEN 0x2808 // Length
72 #define RDH 0x2810 // Head
73 #define RDT 0x2818 // Tail
74
75 // ===== Globals =====
76
77 // Receive Descriptor
78 struct rx_desc
79 {
80     __le64 buffer_addr; // Address of descriptor's data buffer
81     union
82     {
83         __le32 data;
84         struct
85         {
86             __le16 length;
87             __le16 css;
88             }flags;
89     }lower;
90     union
91     {
92         __le32 data;
93         struct
94         {
95             __u8 status;
96             __u8 error;
97             __le16 vlan;
98             }field;
99     }upper;
100 };
101
102 struct buf_info
103 {
104     void* mem;
105     dma_addr_t physical;
106 }buffer_info[MAX];
107
108 // Receive Ring
109 struct rx_ring
110 {
111     void* desc; //Point to ring memory
112     dma_addr_t dma; // Physical addr of ring
113     struct rx_desc * cpu_addr; // descriptor addresses
114     int size; // Length of ring in bytes
115     int count; // Number of desc. in ring
116
117     u16 next_to_use;
118     u16 next_to_clean;
119 };
120
121 // Data that is important will be kept here
122 static struct mydev_dev {
123     // struct net_device netdev;
124     struct cdev cdev;

```

```

125     int input;           // store value for LED register
126     bool status;        // flag for LED, (on or off)
127
128     struct rx_ring* rx_ring;
129     void* hw_addr;       // hold base addr of driver
130     struct work_struct service_task;
131 } mydev;
132
133 // For making device node file
134 static dev_t mydev_node;
135 static struct class *cl;
136
137 //===== Functions =====
138
139 // Workqueue service task
140 // sleep for 0.5seconds, then turn off LEDs
141 static void hw6_service_task(struct work_struct* work)
142 {
143     u32 tail, head;
144
145     head = readl(mydev.hw_addr + RDH);
146     tail = readl(mydev.hw_addr + RDT);
147
148     printk(KERN_INFO "Value of HEAD: %d\n",head);
149     printk(KERN_INFO "Value of TAIL: %d\n",tail);
150
151     printk(KERN_INFO "Service task: SLEEP!\n");
152     msleep(500);
153
154     printk(KERN_INFO "Service task: LEDs off!\n");
155     writel((unsigned int)LED_OFF,mydev.hw_addr + LED_REG);
156
157     // bump tail
158     if(tail >=16)
159         writel(0, mydev.hw_addr + RDT);
160     else
161         writel(tail+1,mydev.hw_addr + RDT);
162
163     // re-enable interrupts
164     writel(IRQ_ENABLE, mydev.hw_addr + IMS);
165 }
166
167 // Interrupt Handler
168 // Turn on both green LEDs, then schedule work.
169 static irqreturn_t hw6_irq_handler(int irq, void* data)
170 {
171     u32 cause;
172     // disable interrupts
173     writel(0xFFFFFFFF, mydev.hw_addr + IMC);
174
175     printk(KERN_INFO "Interrupt: LEDs on!\n");
176     writel((unsigned int)LED_ON, mydev.hw_addr + LED_REG);
177     schedule_work(&mydev.service_task);
178
179     // Read to clear interrupt bit
180     cause = readl(mydev.hw_addr + ICR);
181     printk(KERN_INFO "Cause from ICR: %x\n",cause);
182
183     return IRQ_HANDLED;
184 }
185
186 // Setup resources for ring

```

```

187 static void set_ring(struct pci_dev* pdev)
188 {
189     int i; // for looping
190     unsigned int reg;
191     mydev.rx_ring = kzalloc(sizeof(struct rx_ring), GFP_KERNEL);
192     mydev.rx_ring->count = MAX;
193     mydev.rx_ring->size = sizeof(struct rx_desc) * MAX; //total size of all 16 descriptors
194     mydev.rx_ring->size = ALIGN(mydev.rx_ring->size, 2048);
195     printk(KERN_INFO "ring size set and aligned\n");
196
197     // Allocate and get addresses for ring
198     mydev.rx_ring->desc = dma_alloc_coherent(&pdev->dev, mydev.rx_ring->size, &mydev.
rx_ring->dma, GFP_KERNEL);
199     printk(KERN_INFO "dma_alloc_coherent done!\n");
200
201     reg = (mydev.rx_ring->dma >> 32) & 0xffffffff; // Higher
202     printk(KERN_INFO "Higher: 0x%x \n", reg);
203     writel(reg, mydev.hw_addr+RDBAH);
204
205     reg = (mydev.rx_ring->dma) & 0xffffffff; // Lower
206     printk(KERN_INFO "Lower: 0x%x \n", reg);
207     writel(reg, mydev.hw_addr+RDBAL);
208
209     writel(15, mydev.hw_addr+RDT);
210
211     mydev.rx_ring->next_to_use = 0;
212     mydev.rx_ring->next_to_clean = 0;
213     mydev.rx_ring->cpu_addr = kzalloc(mydev.rx_ring->size, GFP_KERNEL);
214     printk(KERN_INFO "Ring resources set~! Start filling descriptor buffer\n");
215
216     // Set length for Receive Descriptors. Write to RDLEN
217     writel(mydev.rx_ring->size, mydev.hw_addr + RDLEN);
218     for(i=0; i<MAX; i++)
219     {
220         buffer_info[i].mem = kmalloc(2048, GFP_KERNEL);
221         buffer_info[i].physical = dma_map_single(&pdev->dev, buffer_info[i].mem, 2048,
DMA_FROM_DEVICE);
222         mydev.rx_ring->cpu_addr[i].buffer_addr = buffer_info[i].physical;
223     }
224 }
225
226 // Devices supported by this driver
227 static DEFINE_PCI_DEVICE_TABLE(pci_test_tbl) = {
228     { PCI_DEVICE(VENDOR_ID, DEVICE_ID) },
229     { }, /* must have an empty at the end! */
230 };
231
232 // Enable PCI device and map to memory
233 static int my_pci_probe(struct pci_dev *pdev, const struct pci_device_id *ent)
234 {
235     resource_size_t mmio_start, mmio_len;
236     int bars, err;
237
238     /* this is where I'd map BAR's for access, save stuff off, etc. */
239     printk(KERN_INFO "It's dangerous to go alone, take this with you.\n");
240
241     err = pci_enable_device_mem(pdev);
242
243     // set up pci pci connections
244     bars = pci_selectBars(pdev, IORESOURCE_MEM);
245     err = pci_request_selected_regions(pdev, bars, DEVNAME);

```

```

247
248     pci_set_master(pdev);
249
250     // map memory, get base addr of desired device
251     mmio_start = pci_resource_start(pdev, 0);
252     mmio_len = pci_resource_len(pdev, 0);
253     mydev.hw_addr = ioremap(mmio_start, mmio_len);
254
255     // Follow steps in 82583v Controller Datasheet
256     // Steps for Software Initialization under 4.6
257     // disable interrupts
258     writel(0xffffffff, mydev.hw_addr+IMC);
259
260     // Do device reset
261     writel((1 << 26), mydev.hw_addr+CTRL);
262
263     // Modify HEAD to point at zero
264     writel(0, mydev.hw_addr+RDH);
265
266     // Disable Interrupts again
267     writel(0xffffffff, mydev.hw_addr+IMC);
268     readl(mydev.hw_addr+STATUS); // read to flush register
269
270     INIT_WORK(&mydev.service_task, hw6_service_task);
271
272     // Needed for PCIe workarounds - reserved chicken bits
273     // Write GCR bit 22, GCR bit 1
274     writel((readl(mydev.hw_addr+GCR) | (1 << 22), mydev.hw_addr+GCR);
275     writel((readl(mydev.hw_addr+GCR2) | 1, mydev.hw_addr+GCR2);
276
277     // Needed for a forced PHY setup
278     // PHY setup
279     writel(0x1831af08, mydev.hw_addr+MDIC);
280
281     // MAC setup
282     // Set link up while preserving defaults.
283     writel((readl(mydev.hw_addr + CTRL) | 0x40), mydev.hw_addr + CTRL);
284
285     // Clear Status register by reading
286     readl(mydev.hw_addr+STATUS);
287     // Work Queue
288     printk(KERN_INFO "Probe: Work initialized\n");
289
290     // Setup Receive Ring
291     printk(KERN_INFO "Probe: Setup ring resources!\n");
292     set_ring(pdev);
293     printk(KERN_INFO "Probe: Ring Resources set and descriptors filled!\n");
294
295     // Set interrupts. Enable MSI and request IRQ
296     pci_enable_msi(pdev);
297     err = request_irq(pdev->irq, hw6_irq_handler, 0, "Aaron's_IRQ", &mydev);
298
299     // Enable interrupts
300     writel(0x00000000, mydev.hw_addr + IMC); // Not sure this is needed
301     writel(IRQ_ENABLE, mydev.hw_addr + IMS);
302     printk(KERN_INFO "Probe: Interrupts Enabled\n");
303
304     // Enable receiver and setup promiscuous
305     writel(PROMISCUOUS, mydev.hw_addr + RCTL);
306     printk(KERN_INFO "Probe: Receiver enabled!\n");
307
308     /* 0 means success */

```

```

309     return 0;
310
311 }
312
313 // Clean up PCI allocations, disable device
314 static void my_pci_remove(struct pci_dev *pdev)
315 {
316     int i;
317
318     // Cleanup Work Queue
319     cancel_work_sync(&mydev.service_task);
320
321     // Disable interrupts
322     free_irq(pdev->irq, &mydev);
323     pci_disable_msi(pdev);
324
325     // Free and unpin memory for buffer info
326     for(i=0; i<MAX; i++)
327     {
328         kfree(buffer_info[i].mem);
329         dma_unmap_single(&pdev->dev, buffer_info->physical, 2048, DMA_TO_DEVICE);
330     }
331
332     // Free ring
333     dma_free_coherent(&pdev->dev, mydev.rx_ring->size, mydev.rx_ring->desc, mydev.rx_ring->dma);
334     kfree(mydev.rx_ring->cpu_addr);
335     kfree(mydev.rx_ring);
336
337     // unmap pci device
338     iounmap(mydev.hw_addr);
339     pci_release_selected_regions(pdev, pci_select_bars(pdev, IORESOURCE_MEM));
340     pci_disable_device(pdev);
341
342     printk(KERN_INFO "So long!!\n");
343 }
344
345 // Name of my driver and associated functions
346 static struct pci_driver my_pci_driver = {
347     .name = DEVNAME,
348     .id_table = pci_test_tbl,
349     .probe = my_pci_probe,
350     .remove = my_pci_remove,
351 };
352
353 // Open function
354 static int pci_hw6_open(struct inode *inode, struct file *file)
355 {
356     printk(KERN_INFO "(my_pci_driver) successfully opened!\n");
357     return 0;
358 }
359
360 // Release function
361 static int pci_hw6_release(struct inode *inode, struct file *file)
362 {
363     printk(KERN_INFO "(my_pci_driver) successfully closed!\n");
364     return 0;
365 }
366
367 // Read function
368 static ssize_t pci_hw6_read(struct file *file, char __user *buf,
369                             size_t len, loff_t *offset)
370 {

```

```

370     /* Get a local kernel buffer set aside */
371     int ret;
372     u32 head_tail;
373     head_tail= (readl(mydev.hw_addr + RDH) << 16) | readl(mydev.hw_addr + RDT);
374
375     if (*offset >= sizeof(int))
376         return 0;
377
378     /* Make sure our user wasn't bad... */
379     if (!buf) {
380         ret = -EINVAL;
381         goto out;
382     }
383
384     // Pass blink rate value to userspace
385     if (copy_to_user(buf, &head_tail, sizeof(unsigned int))) {
386         ret = -EFAULT;
387         goto out;
388     }
389     ret = sizeof(unsigned int);
390     *offset += len;
391
392     /* Good to go, so printk the thingy */
393     printk(KERN_INFO "(my_pci_driver:read)User got from us %d\n",head_tail);
394
395 out:
396     return ret;
397 }
398
399 // Write function
400 static ssize_t pci_hw6_write(struct file *file, const char __user *buf,
401                             size_t len, loff_t *offset)
402 {
403     int ret;
404     /* Make sure our user isn't bad... */
405     if (!buf) {
406         ret = -EINVAL;
407         goto out;
408     }
409
410     /* Copy from the user-provided buffer */
411     if (copy_from_user(&mydev.input, buf, len)) {
412         /* uh-oh... */
413         ret = -EFAULT;
414         goto out;
415     }
416
417     if(mydev.input < 0)
418     {
419         printk("(my_pci_driver:write)User wrote negative value. Return error\n");
420         ret = EINVAL;
421         goto out;
422     }
423     else if(mydev.input == 0)
424         printk("(my_pci_driver:write)User wrote 0. Do nothing\n");
425     else
426     {
427         printk("(my_pci_driver:write)User wrote %d\n",mydev.input);
428         blink_rate = mydev.input;
429     }
430     ret = len;
431

```

```
432 out:
433     return ret;
434 }
435
436 /* File operations for our device */
437 static struct file_operations mydev_fops = {
438     .owner = THIS_MODULE,
439     .open = pci_hw6_open,
440     .read = pci_hw6_read,
441     .write = pci_hw6_write,
442     .release = pci_hw6_release,
443 };
444
445 // Initialization
446 static int __init pci_hw6_init(void)
447 {
448     mydev.status = false;
449
450     printk(KERN_INFO "(my_pci_driver) module loading...\n");
451
452     if (alloc_chrdev_region(&mydev_node, 0, DEVCNT, DEVNAME)) {
453         printk(KERN_ERR "alloc_chrdev_region() failed!\n");
454         return -1;
455     }
456
457     // Get major number for device
458     printk(KERN_INFO "Allocated %d devices at major: %d\n", DEVCNT,
459            MAJOR(mydev_node));
460
461     // Create node file. No need for mknod
462     if((cl = class_create( THIS_MODULE, DEVNAME)) == NULL)
463     {
464         printk(KERN_ALERT "Class creation failed\n");
465         unregister_chrdev_region(mydev_node, DEVCNT);
466         return -1;
467     }
468     if(device_create(cl, NULL, mydev_node, NULL, DEV_NODE_NAME) == NULL)
469     {
470         printk(KERN_ALERT "Device creation failed\n");
471         class_destroy(cl);
472         unregister_chrdev_region(mydev_node, DEVCNT);
473     }
474
475     /* Initialize the character device and add it to the kernel */
476     cdev_init(&mydev.cdev, &mydev_fops);
477     mydev.cdev.owner = THIS_MODULE;
478
479     if (cdev_add(&mydev.cdev, mydev_node, DEVCNT)) {
480         printk(KERN_ERR "cdev_add() failed!\n");
481         /* clean up chrdev allocation */
482         unregister_chrdev_region(mydev_node, DEVCNT);
483
484         return -1;
485     }
486     printk(KERN_INFO "Node created\n");
487
488     printk(KERN_INFO "(my_pci_driver) Registering PCI Driver...\n");
489     return pci_register_driver(&my_pci_driver);
490 }
491
492 // Clean up when removing driver
493 static void __exit pci_hw6_exit(void)
```



```
494 {
495     // Disable interrupts
496     writel(0xffffffff, mydev.hw_addr+IMC);
497
498     /* destroy the cdev */
499     cdev_del(&mydev.cdev);
500     device_destroy(cl, mydev_node);
501     class_destroy(cl);
502
503     /* Unregister PCI Driver*/
504     pci_unregister_driver(&my_pci_driver);
505
506     /* clean up the devices */
507     unregister_chrdev_region(mydev_node, DEVCNT);
508     printk(KERN_INFO "(my_pci_driver) module unloaded!\n");
509
510 }
511
512 MODULE_AUTHOR("Aaron Chan");
513 MODULE_LICENSE("GPL");
514 MODULE_VERSION("0.2");
515 module_init(pci_hw6_init);
516 module_exit(pci_hw6_exit);
517
```